

# titanic (2)

July 2, 2024

## 1 Titanic Data Analysis and Classification

### 1.0.1 Imports and Data Preparation

```
[1]: %%capture
!pip install tensorflow;
```

```
[2]: import numpy as np
import pandas as pd
from sklearn.feature_selection import chi2
from sklearn.feature_selection import f_classif
from sklearn import tree
import re
import math
from tensorflow import keras
from sklearn.preprocessing import MinMaxScaler
```

```
[3]: train = pd.read_csv('data/train.csv')
test = pd.read_csv('data/test.csv')
```

```
[4]: data = pd.concat([train, test])
```

```
[5]: data[7:10]
```

```
[5]: PassengerId  Survived  Pclass  \
7              8         0.0        3
8              9         1.0        3
9             10         1.0        2
```

```

              Name      Sex  Age  SibSp  \
7              Palsson, Master. Gosta Leonard   male    2.0         3
8  Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg) female   27.0         0
9              Nasser, Mrs. Nicholas (Adele Achem) female   14.0         1
```

```

Parch  Ticket     Fare  Cabin  Embarked
7      1  349909   21.0750   NaN         S
8      2  347742   11.1333   NaN         S
```

9      0   237736   30.0708   NaN      C

```
[6]: data['InTrain'] = data['PassengerId'].apply(lambda x: True if x <= len(train)
↳ else False)
```

## 1.0.2 Ticket Class

### Summary


```
[7]: data['Pclass'].describe()
```

```
[7]: count      1309.000000
     mean         2.294882
     std         0.837836
     min         1.000000
     25%         2.000000
     50%         3.000000
     75%         3.000000
     max         3.000000
     Name: Pclass, dtype: float64
```

### Visualization

```
[8]: data['Pclass'].value_counts().plot.pie()
```

```
[8]: <AxesSubplot:ylabel='Pclass'>
```



### Correlation

```
[9]: corr = chi2(pd.DataFrame(data['Pclass'][0:len(train)]), data['Survived'][0:
    ↪len(train)])
    print(f'Chi2 = {corr[0]}\nP-value = {corr[1]}')
```

Chi2 = [30.87369944]

P-value = [2.75378563e-08]

### 1.0.3 Sex

#### Summary


```
[10]: data['Sex'].describe()
```

```
[10]: count      1309
      unique        2
      top      male
      freq      843
      Name: Sex, dtype: object
```

#### Visualization

```
[11]: data['Sex'].value_counts().plot.pie()
```

```
[11]: <AxesSubplot:ylabel='Sex'>
```



#### Correlation

```
[12]: data['Sex'] = data['Sex'].apply(lambda x: 1 if x == 'male' else 0)
```

```
corr = chi2(pd.DataFrame(data['Sex'][0:len(train)]), data['Survived'][0:
    ↪len(train)])
print(f'Chi2 = {corr[0]}\nP-value = {corr[1]}')
```

Chi2 = [92.70244698]

P-value = [6.07783826e-22]

#### 1.0.4 Fare

##### Summary

```
[13]: data['Fare'].describe()
```

```
[13]: count    1308.000000
      mean      33.295479
      std      51.758668
      min       0.000000
      25%       7.895800
      50%      14.454200
      75%      31.275000
      max      512.329200
      Name: Fare, dtype: float64
```


##### Imputation

```
[14]: fare_avg = data['Fare'].mean()
      data['Fare'].fillna(fare_avg, inplace = True);
```

##### Visualization

```
[15]: data['Fare'].plot.hist(bins = 50)
```

```
[15]: <AxesSubplot:ylabel='Frequency'>
```



### Correlation

```
[16]: corr = f_classif(pd.DataFrame(data['Fare'][0:len(train)]), data['Survived'][0:
    ↪len(train)])
print(f'ANOVA = {corr[0]}\nP-value = {corr[1]}')
```

ANOVA = [63.03076423]

P-value = [6.12018934e-15]

### 1.0.5 Age / Age Group

#### Feature Engineering

```
[17]: def assignAgeGroup(x):
    if x < 18:
        return 0
    elif x < 65:
        return 1
    elif x >= 65:
        return 2
    return 3

data['Age Group'] = data['Age'].apply(assignAgeGroup)
```

### Summary

```
[18]: data['Age'].describe()
```

```
[18]: count    1046.000000
      mean      29.881138
      std       14.413493
      min        0.170000
      25%       21.000000
      50%       28.000000
      75%       39.000000
      max       80.000000
      Name: Age, dtype: float64
```


```
[19]: data['Age Group'].describe()
```

```
[19]: count    1309.000000
      mean      1.294118
      std       0.919449
      min        0.000000
      25%        1.000000
      50%        1.000000
      75%        1.000000
      max        3.000000
      Name: Age Group, dtype: float64
```

### Visualization


```
[20]: data['Age'].plot.hist(bins = 20)
```

```
[20]: <AxesSubplot:ylabel='Frequency'>
```



```
[21]: data['Age Group'].value_counts().plot.pie()
```

```
[21]: <AxesSubplot:ylabel='Age Group'>
```



### Correlation

```
[22]: corr = f_classif(pd.DataFrame(data['Age'][0:len(train)]).dropna(),  
    ↪data['Survived'][0:len(train)][data['Age'][0:len(train)].notna()])  
print(f'Age:\nANOVA = {corr[0]}\nP-value = {corr[1]}')
```

Age:

ANOVA = [4.27119493]

P-value = [0.03912465]

```
[23]: corr = chi2(pd.DataFrame(data['Age Group'][0:len(train)]), data['Survived'][0:  
    ↪len(train)])  
print(f'Age Group:\nChi2 = {corr[0]}\nP-value = {corr[1]}')
```

Age Group:

Chi2 = [10.28445646]

P-value = [0.00134156]

### 1.0.6 Cabin

#### Summary

```
[24]: data['Cabin'].describe()
```

```
[24]: count          295
      unique         186
      top      C23 C25 C27
      freq           6
      Name: Cabin, dtype: object
```

### Imputation

```
[25]: data['Last Name'] = data['Name'].apply(lambda x: x.split()[0][0:-1])
```

```
[26]: lastnames = data['Last Name'].unique()
      cabins = ['U'] * len(data)

      for lastname in lastnames:
          family = data[data['Last Name'] == lastname]
          if len(family['Cabin'].dropna()) != 0:
              cabin = family['Cabin'].dropna().tolist()[0]
              pIds = family['PassengerId'].tolist()
              for pId in pIds:
                  cabins[pId-1] = cabin

      data['Cabin'] = cabins
```

```
[27]: data['Cabin'].describe()
```

```
[27]: count          1309
      unique          172
      top             U
      freq           977
      Name: Cabin, dtype: object
```

### Feature Engineering and Imputation

```
[28]: data['Floor'] = data['Cabin'].apply(lambda x: x[0])
```

```
[29]: room_sum = 0
      for cabin in cabins:
          if len(cabin) != 1:
              room_sum += float(re.findall(r'\d+', cabin[-2:])[0])
      room_avg = room_sum / 332

      data['Room'] = data['Cabin'].apply(lambda x: room_avg if len(x) == 1 else
      ↪ float(re.findall(r'\d+', x[-2:])[0]))
```

```
[30]: data['Room'].describe()
```




```
[30]: count    1309.000000
      mean      37.311810
      std       13.499729
      min        1.000000
      25%       37.141566
      50%       37.141566
      75%       37.141566
      max       99.000000
      Name: Room, dtype: float64
```

### Visualization


```
[31]: data['Floor'][data['Floor'] != 3.5].value_counts().plot.pie()
```

```
[31]: <AxesSubplot:ylabel='Floor'>
```



```
[32]: data['Room'][data['Room'] != room_avg].plot.hist(bins = 20)
```

```
[32]: <AxesSubplot:ylabel='Frequency'>
```



### Correlation

```
[33]: data['Floor'] = data['Floor'].replace(['A', 'B', 'C', 'D', 'E', 'F', 'G', 'T', 'U'],
      ↪ ['U'], [0, 1, 2, 3, 4, 5, 6, 7, 3.5]);
corr = f_classif(pd.DataFrame(data['Floor'][0:len(train)]), data['Survived'][0:
      ↪ len(train)])
print(f'Floor:\nANOVA = {corr[0]}\nP-value = {corr[1]}')
```

Floor:

ANOVA = [22.1482981]

P-value = [2.92578455e-06]

```
[34]: corr = f_classif(pd.DataFrame(data['Room'][0:len(train)] [data['Room'][0:
      ↪ len(train)] != room_avg]), data['Survived'][0:len(train)] [data['Room'][0:
      ↪ len(train)] != room_avg])
print(f'Room:\nANOVA = {corr[0]}\nP-value = {corr[1]}')
```

Room:

ANOVA = [0.50850824]

P-value = [0.47652756]

### 1.0.7 Family Size

#### Feature Engineering

```
[35]: family_size = [None] * len(data)
```

```

for lastname in lastnames:
    family = data[data['Last Name'] == lastname]
    pIds = family['PassengerId'].tolist()
    for pId in pIds:
        family_size[pId-1] = len(family)


data['Family Size'] = family_size

```

## Visualization

```
[36]: data['Family Size'].value_counts().plot.pie()
```

```
[36]: <AxesSubplot:ylabel='Family Size'>
```



## Correlation

```

[37]: corr = f_classif(pd.DataFrame(data['Family Size'][0:len(train)]),
    ↪ data['Survived'][0:len(train)])
print(f'Family Size:\nANOVA = {corr[0]}\nP-value = {corr[1]}')

```

```

Family Size:
ANOVA = [1.06329472]
P-value = [0.30274538]

```

## 1.0.8 Family Survived

### Feature Engineering

```
[38]: family_survived = [None] * len(data)


for lastname in lastnames:
    family = data[data['Last Name'] == lastname]
    survived = family['Survived'].tolist()
    avg_survived = 0
    for member in survived:
        if math.isnan(member):
            avg_survived += 0.5
        else:
            avg_survived += member
    avg_survived /= len(family)
    pIds = family['PassengerId'].tolist()
    for pId in pIds:
        family_survived[pId-1] = avg_survived

data['Family Survived'] = family_survived
```

### Visualization

```
[39]: data['Family Survived'].plot.hist(bins = 10)
```

```
[39]: <AxesSubplot:ylabel='Frequency'>
```



### Correlation

```
[40]: corr = f_classif(pd.DataFrame(data['Family Survived'])[0:len(train)],
    ↪data['Survived'])[0:len(train))
print(f'Family Survived:\nANOVA = {corr[0]}\nP-value = {corr[1]}')
```

Family Survived:  
ANOVA = [3197.30644761]  
P-value = [1.07334179e-296]


### 1.0.9 Model and Classification

#### Decision Tree

```
[41]: final_train = data[['Pclass', 'Sex', 'Age Group', 'Floor', 'Family Survived']][:
    ↪len(train)]
final_test = data[['Pclass', 'Sex', 'Age Group', 'Floor', 'Family
    ↪Survived']][len(train):]
```

```
[42]: d_tree = tree.DecisionTreeClassifier()
d_tree.fit(final_train, data['Survived'][:len(final_train)]);
```

```
[43]: tree.plot_tree(d_tree);
```



```
[44]: classifications = d_tree.predict(final_test)
final = pd.DataFrame()
final['Survived'] = classifications
final['PassengerId'] = data['PassengerId'][len(train):]
final.to_csv('submission.csv');
```

## Neural Network

```
[45]: # Standardize dataset
      scaler = MinMaxScaler()
      final_train = scaler.fit_transform(final_train)
      final_test = scaler.transform(final_test)

[46]: model = keras.Sequential()
      model.add(keras.layers.Dense(128, activation = 'sigmoid'))
      model.add(keras.layers.Dense(128, activation = 'sigmoid'))
      model.add(keras.layers.Dense(1, activation = 'sigmoid'))

[47]: model.compile(loss = 'binary_crossentropy', optimizer = 'adam', metrics = [
      ↪['accuracy'])

[52]: model.fit(x = final_train, y = data['Survived'][:len(final_train)], epochs = 15)
```

```
Epoch 1/15
28/28 [=====] - 0s 963us/step - loss: 0.2088 -
accuracy: 0.9371
Epoch 2/15
28/28 [=====] - 0s 889us/step - loss: 0.1904 -
accuracy: 0.9394
Epoch 3/15
28/28 [=====] - 0s 778us/step - loss: 0.1730 -
accuracy: 0.9484
Epoch 4/15
28/28 [=====] - 0s 815us/step - loss: 0.1628 -
accuracy: 0.9450
Epoch 5/15
28/28 [=====] - 0s 778us/step - loss: 0.1493 -
accuracy: 0.9585
Epoch 6/15
28/28 [=====] - 0s 852us/step - loss: 0.1422 -
accuracy: 0.9562
Epoch 7/15
28/28 [=====] - 0s 778us/step - loss: 0.1348 -
accuracy: 0.9562
Epoch 8/15
28/28 [=====] - 0s 778us/step - loss: 0.1340 -
accuracy: 0.9607
Epoch 9/15
28/28 [=====] - 0s 815us/step - loss: 0.1227 -
accuracy: 0.9630
Epoch 10/15
28/28 [=====] - 0s 778us/step - loss: 0.1218 -
accuracy: 0.9574
Epoch 11/15
```

```

28/28 [=====] - 0s 797us/step - loss: 0.1222 -
accuracy: 0.9562
Epoch 12/15
28/28 [=====] - 0s 796us/step - loss: 0.1143 -
accuracy: 0.9607
Epoch 13/15
28/28 [=====] - 0s 778us/step - loss: 0.1126 -
accuracy: 0.9618
Epoch 14/15
28/28 [=====] - 0s 748us/step - loss: 0.1083 -
accuracy: 0.9663
Epoch 15/15
28/28 [=====] - 0s 741us/step - loss: 0.1108 -
accuracy: 0.9574

```

```
[52]: <keras.callbacks.History at 0x2466d269d30>
```

```
[53]: predictions = model.predict(final_test)
```

```
14/14 [=====] - 0s 696us/step
```

```
[54]: rounded_predictions = [round(p[0]) for p in predictions.tolist()]
```

```
[55]: final2 = pd.DataFrame()
final2['Survived'] = rounded_predictions
final2['PassengerId'] = data['PassengerId'][len(train):]
final2.to_csv('submission2.csv');
```

### 1.0.10 Results

Kaggle verified accuracy: 78.71%

Leaderboard position: 1775 / 15841 as of April 1st, 2024